

用户调用Nova API挂载磁盘

```
1 os-volume_attachments
2 {
3     "volumeAttachment": {
4         "device": "/dev/vdb",
5         "serverId": "6edad723-c62d-40d4-a29e-d9ac6a3247ee",
6         "id": "3b9fdb6c-200e-494d-aa71-f47c70e87d7e",
7         "volumeId": "3b9fdb6c-200e-494d-aa71-f47c70e87d7e"
8     }
9 }
```

```
1 API接口相应nova volume-attach请求:
2 nova.api.openstack.compute.volumes.VolumeAttachmentController
3 def create(self, req, server_id, body)
```

```
1 nova.compute.api.API
2     @check_instance_lock
3     @check_instance_state(vm_state=[vm_states.ACTIVE, vm_states.PAUSED, vm_states.STOPPED, vm_states.RESIZED,
4                                     vm_states.SOFT_DELETED, vm_states.SHELVED,
5                                     vm_states.SHELVED_OFFLOADED])
6     def attach_volume(self, context, instance, volume_id, device=None,
7                       disk_bus=None, device_type=None):
8         """Attach an existing volume to an existing instance."""
9         if device and not block_device.match_device(device):
10             raise exception.InvalidDevicePath(path=device)
11
12         is_shelved_offloaded = instance.vm_state == vm_states.SHELVED_OFFLOADED
13         if is_shelved_offloaded:
14             """
15             shelved_offloaded状态下的虚拟机也允许挂载,shelved_offloaded状态下虚拟机从hypervisor层面清除,
16             在_attach_volume_shelved_offloaded方法中虚拟机并没有实际在KVM层面执行挂载,只是创建block_device_mapping表的记录
17             并且通知Cinder修改卷状态为in-use,并且直接创建attachment记录
18             """
19             return self._attach_volume_shelved_offloaded(context, instance, volume_id,
20                                                         device, disk_bus, device_type)
21
22         return self._attach_volume(context, instance, volume_id, device,
23                                    disk_bus, device_type)
```

```
1 nova.compute.api.API
2     def _attach_volume(self, context, instance, volume_id, device,
3                        disk_bus, device_type):
4         """Attach an existing volume to an existing instance.
5
6         This method is separated to make it possible for cells version
7         to override it.
```

```

8      """
9      """
10     _create_volume_bdm方法负责在block_device_mapping表中创建一个记录，接口会通过API节点创建，
11     或者API节点调用nova_compute节点的RPC接口reserve_block_device_name来创建！
12     """
13     volume_bdm = self._create_volume_bdm(
14         context, instance, device, volume_id, disk_bus=disk_bus,
15         device_type=device_type)
16     try:
17         """
18         可用域检查与卷状态置attaching
19         """
20         self._check_attach_and_reserve_volume(context, volume_id, instance)
21         self.compute_rpcapi.attach_volume(context, instance, volume_bdm)
22     except Exception:
23         with excutils.save_and_reraise_exception():
24             volume_bdm.destroy()
25
26     return volume_bdm.device_name
27

```

```

1 nova.compute.manager.ComputeManager
2     @wrap_exception()
3     @wrap_instance_fault
4     def attach_volume(self, context, instance, bdm):
5         """Attach a volume to an instance."""
6         """
7         nova.virt.block_device --> driver_block_device
8         convert_volume方法的主要作用是主要作用是将Block_Device_Mapping中的ORM对象转换成
9         DriverVolumeBlockDevice、DriverSnapshotBlockDevice、DriverImageBlockDevice、DriverBlankBlockDevice等对象
10        """
11        driver_bdm = driver_block_device.convert_volume(bdm)
12
13        """
14        对下面的代码用uuid加锁
15        """
16        @utils.synchronized(instance.uuid)
17        def do_attach_volume(context, instance, driver_bdm):
18            try:
19                return self._attach_volume(context, instance, driver_bdm)
20            except Exception:
21                with excutils.save_and_reraise_exception():
22                    bdm.destroy()
23
24        do_attach_volume(context, instance, driver_bdm)

```

```

1 nova.compute.manager.ComputeManager
2     def _attach_volume(self, context, instance, bdm):
3         context = context.elevated()
4         LOG.info(_LI('Attaching volume %(volume_id)s to %(mountpoint)s'),

```

```

5         {'volume_id': bdm.volume_id,
6          'mountpoint': bdm['mount_device']},
7         instance=instance)
8     try:
9         """
10        DriverVolumeBlockDevice、DriverSnapshotBlockDevice、DriverImageBlockDevice、DriverBlankBlockDevice都有attach方法
11        但是他们的接口稍有不同，除了DriverVolumeBlockDevice，其他三种类型都需要在cinder处创建卷，因此需要传递wait_func参数
12        换句话说DriverSnapshotBlockDevice、DriverImageBlockDevice、DriverBlankBlockDevice的attach包含的volume_create
13        主要用于创建虚机的流程。
14        如果虚机创建完成，将一个已有卷挂载到VM，走DriverVolumeBlockDevice的attach接口！
15        """
16        bdm.attach(context, instance, self.volume_api, self.driver,
17                   do_check_attach=False, do_driver_attach=True)
18    except Exception:
19        with excutils.save_and_reraise_exception():
20            LOG.exception(_LE("Failed to attach %(volume_id)s "
21                              "at %(mountpoint)s"),
22                          {'volume_id': bdm.volume_id,
23                           'mountpoint': bdm['mount_device']},
24                          instance=instance)
25            self.volume_api.unreserve_volume(context, bdm.volume_id)
26
27    info = {'volume_id': bdm.volume_id}
28    self._notify_about_instance_usage(
29        context, instance, "volume.attach", extra_usage_info=info)

```

nova.virt.block_device.DriverVolumeBlockDevice中的attach方法：

整个过程如下：

1. 调用os-brick获取本机信息，以及传输协议需要的信息
---->nova.virt.block_device.DriverVolumeBlockDevice.attach
2. 通过cinderAPI获取VOLUME的连接信息
---->nova.virt.block_device.DriverVolumeBlockDevice.attach
3. 调用协议驱动建立连接，如NFS协议mount关键目录，iscsi协议调用os-brick建立连接等（该过程发生在nova_compute容器中）！
--->nova.virt.driver.LibvirtDriver.attach_volume
4. nova-compute通知nova-libvirt创建virtio设备驱动，并更新虚机的XML配置(该过程发生在nova_libvirt容器中)！
--->nova.virt.driver.LibvirtDriver.attach_volume
5. 通知Cinder修改状态，并且完成数据库的最后更新！
---->nova.virt.block_device.DriverVolumeBlockDevice.attach

```

1 nova.virt.block_device.DriverVolumeBlockDevice
2     @update_db
3     def attach(self, context, instance, volume_api, virt_driver,
4               do_check_attach=True, do_driver_attach=False, **kwargs):
5
6         volume = volume_api.get(context, self.volume_id)
7         if do_check_attach:
8             """
9             检查VOLUME的状态，挂载虚机的时候是不检查的，考虑到multi-attach??
10            """

```

```

11         volume_api.check_attach(context, volume, instance=instance)
12
13     volume_id = volume['id']
14     context = context.elevated()
15     """
16     libvirt.LibvirtDriver --> virt_driver
17     os_brick.initiator.connector, 收集本机的信息, 以及一些存储协议建立连接时所用的信息。
18     """
19     connector = virt_driver.get_volume_connector(instance)
20     """
21     向Cinder请求卷的连接信息 --> os_initialize_connection
22     """
23     connection_info = volume_api.initialize_connection(context, volume_id, connector)
24     if 'serial' not in connection_info:
25         connection_info['serial'] = self.volume_id
26     self._preserve_multipath_id(connection_info)
27
28     """
29     对于active状态下的VM进行attach, do_driver_attach必定为true;
30     而创建过程中一些卷的attach, do_driver_attach为false, 应在虚拟机启动过程中, 已经执行attach了
31     """
32
33     if do_driver_attach:
34         encryption = encryptors.get_encryption_metadata(
35             context, volume_api, volume_id, connection_info)
36
37         try:
38             """
39             Libvirt驱动中的attach方法建立连接!
40
41             """
42
43             virt_driver.attach_volume(context, connection_info, instance,
44                                     self['mount_device'], disk_bus=self['disk_bus'],
45                                     device_type=self['device_type'], encryption=encryption)
46
47         except Exception:
48             with excutils.save_and_reraise_exception():
49                 LOG.exception(_LE("Driver failed to attach volume "
50                                "%(volume_id)s at %(mountpoint)s"),
51                               {'volume_id': volume_id,
52                                'mountpoint': self['mount_device']},
53                               instance=instance)
54                 volume_api.terminate_connection(context, volume_id,
55                                                  connector)
56
57     self['connection_info'] = connection_info
58     if self.volume_size is None:
59         self.volume_size = volume.get('size')
60
61     mode = 'rw'
62     if 'data' in connection_info:
63         mode = connection_info['data'].get('access_mode', 'rw')
64
65     if volume['multiattach'] or volume['attach_status'] == "detached":
66         # NOTE(mriedem): save our current state so connection_info is in
67         # the database before the volume status goes to 'in-use' because
68         # after that we can detach and connection_info is required for

```

```

69         # detach.
70         """
71         更新Block_Device_Mapping表中的信息
72         """
73         self.save()
74         try:
75             """
76             cinder的attach接口主要工作是更新volume的状态，已经attachments表中的记录！
77             并且该接口为driver提供了一个回调方法，BaseVD.attach()该方法没有实际操作，可以用来定制！
78             """
79             volume_api.attach(context, volume_id, instance.uuid, self['mount_device'], mode=mode)
80         except Exception:
81             with excutils.save_and_reraise_exception():
82                 if do_driver_attach:
83                     try:
84                         virt_driver.detach_volume(connection_info,
85                                                    instance,
86                                                    self['mount_device'],
87                                                    encryption=encryption)
88                     except Exception:
89                         LOG.warning(_LW("Driver failed to detach volume "
90                                         "%(volume_id)s at %(mount_point)s."),
91                                   {'volume_id': volume_id,
92                                    'mount_point': self['mount_device']},
93                                   exc_info=True, instance=instance)
94                         volume_api.terminate_connection(context, volume_id,
95                                                         connector)
96
97             # Cinder-volume might have completed volume attach. So
98             # we should detach the volume. If the attach did not
99             # happen, the detach request will be ignored.
100            volume_api.detach(context, volume_id)


```

关于Libvirt实现的不同存储协议：

```

1  #包路径: nova.virt.libvirt.volume
2  @profiler.trace_cls("volume_api")
3  class LibvirtBaseVolumeDriver(object):
4      """Base class for volume drivers."""
5      def __init__(self, host, is_block_dev):
6          self.host = host
7          self.is_block_dev = is_block_dev
8
9      def get_config(self, connection_info, disk_info):
10         """Returns xml for libvirt."""
11         """生成该存储在VM中使用的XML"""
12         return conf
13
14     def connect_volume(self, connection_info, disk_info):
15         """Connect the volume."""
16         pass
17
18     def disconnect_volume(self, connection_info, disk_dev):
19         """Disconnect the volume."""
20         pass

```



cinder的ATTACH接口，这个接口实际上是NOVA完成ATTACH操作以后的一个回调接口，主要用途是修改状态，以及创建attachment记录！

```
1 cinder.volume.manager.VolumeManager
2 @coordination.synchronized('{volume_id}')
3 def attach_volume(self, context, volume_id, instance_uuid, host_name,
4                   mountpoint, mode, volume=None):
5     """Updates db to show volume is attached."""
6     # FIXME(lixiaoy1): Remove this in v4.0 of RPC API.
7     if volume is None:
8         # For older clients, mimic the old behavior and look
9         # up the volume by its volume_id.
10        volume = objects.Volume.get_by_id(context, volume_id)
11        # Get admin_metadata. This needs admin context.
12        with volume.obj_as_admin():
13            volume_metadata = volume.admin_metadata
14        # check the volume status before attaching
15        if volume.status == 'attaching':
16            if (volume_metadata.get('attached_mode') and volume_metadata.get('attached_mode') != mode):
17                raise exception.InvalidVolume(reason=_("being attached by different mode"))
18
19        if (volume.status == 'in-use' and not volume.multiattach and not volume.migration_status):
20            raise exception.InvalidVolume(reason=_("volume is already attached"))
21
22        host_name_sanitized = utils.sanitize_hostname(
23            host_name) if host_name else None
24        if instance_uuid:
25            attachments = (VA_LIST.get_all_by_instance_uuid(context, instance_uuid))
26        else:
27            attachments = (VA_LIST.get_all_by_host(context, host_name_sanitized))
28        if attachments:
29            # check if volume->instance mapping is already tracked in DB
30            """
31            检查attachments表的记录如果记录已经存在，将volume.status改为in-use，并且返回相应attachment!
32            """
```

```

33         for attachment in attachments:
34             if attachment['volume_id'] == volume_id:
35                 volume.status = 'in-use'
36                 volume.save()
37                 return attachment
38
39     self._notify_about_volume_usage(context, volume, "attach.start")
40     """
41     创建attachment表记录, 将attachment状态改为ATTACHING
42     """
43     attachment = volume.begin_attach(mode)
44
45     if instance_uuid and not uuidutils.is_uuid_like(instance_uuid):
46         attachment.attach_status = (fields.VolumeAttachStatus.ERROR_ATTACHING)
47         attachment.save()
48         raise exception.InvalidUUID(uuid=instance_uuid)
49
50     if volume_metadata.get('readonly') == 'True' and mode != 'ro':
51         attachment.attach_status = (fields.VolumeAttachStatus.ERROR_ATTACHING)
52         attachment.save()
53         self.message_api.create(
54             context, defined_messages.EventIds.ATTACH_READONLY_VOLUME,
55             context.project_id, resource_type=resource_types.VOLUME,
56             resource_uuid=volume.id)
57         raise exception.InvalidVolumeAttachMode(mode=mode, volume_id=volume.id)
58
59     try:
60         utils.require_driver_initialized(self.driver)
61
62         LOG.info(_LI('Attaching volume %(volume_id)s to instance '
63                     '%(instance)s at mountpoint %(mount)s on host '
64                     '%(host)s.'),
65                  {'volume_id': volume_id, 'instance': instance_uuid,
66                   'mount': mountpoint, 'host': host_name_sanitized},
67                  resource=volume)
68
69         """
70         关于驱动的一个回调函数! 可以为一些定制预留接口!
71         """
72         self.driver.attach_volume(context,
73                                   volume,
74                                   instance_uuid,
75                                   host_name_sanitized,
76                                   mountpoint)
77     except Exception:
78         with excutils.save_and_reraise_exception():
79             attachment.attach_status = (
80                 fields.VolumeAttachStatus.ERROR_ATTACHING)
81             attachment.save()
82         """
83         完成attach
84         """
85     volume = attachment.finish_attach(
86         instance_uuid,
87         host_name_sanitized,
88         mountpoint,
89         mode)
90

```

```
91     self._notify_about_volume_usage(context, volume, "attach.end")
92     LOG.info(_LI("Attach volume completed successfully."),
93             resource=volume)
94     return attachment
```